CS 511: Assignment #2
Using Semaphores for Buffer Synchronization

Due Wednesday, October 16 at 11pm.

1   Assignment

The purpose of this assignment is to give you some experience writing a simple
multi-threaded program whose threads are synchronized by use of a semaphore.

In this assignment one thread will read lines of text from a file then write the lines
into a buffer in memory. A second thread running simultaneously will copy the
lines from the buffer then write them to an output file. When the program finishes,
the output file should be equal to the input file. The concurrency problem is to
synchronize the threads's access to the buffer.

1.1   Part 1

As a first step, write a single-threaded program named "rw" that repeatedly reads
a line of text from an input file, writes the line into a buffer, then writes the line
from the buffer to the output file. The program should take two command line
arguments, the first being the name of the input file, the second being the name of
the output file.

The purpose of this part of the assignment is to have you write and debug the file-
handling code before tackling the concurrency aspects. Library calls that might
be useful include `fopen(3)`, `getline(3)`, and `fwrite(3)`,

1.2   Part 2

Write a two-threaded program named "transfer1" in which one thread (the "fill
thread") reads lines of text from the input file and writes them into a buffer. A sec-
ond thread running simultaneously (the "drain thread") copies the lines from the
buffer then writes them to an output file. You must synchronize the two threads's

access to the buffer using a single semaphore. The buffer should be initialized to all zeroes.

Once the fill thread has detected end of the input file it should write "QUIT" into the buffer. When the drain thread reads QUIT it should terminate. The fill thread should wait for the drain thread to terminate and, after the drain thread has terminated, the fill thread should deallocate the buffer and the semaphore then terminate the program.

The fill thread should print a message whenever it writes a line into the buffer. The drain thread should print a message whenever it reads a line from the buffer.

The fill thread should write each line into the buffer following the previous line. You need not worry about the buffer overflowing – allocate a buffer that is big enough to hold all the lines of the input file, including the terminating null character for each string. The drain thread must be able to examine the buffer and determine whether a new line has been written into it since the last time the drain thread executed. To make this determination, the drain thread should check whether the next byte it plans to read still has the initial zero value or something different; if the byte is no longer zero that means a string has been written into the buffer starting at that position.

The drain and fill threads will run at unpredictable rates, based on the operating system's scheduling decisions. Your program should operate properly regardless how often or when either thread executes. So that we can explore different thread execution rates, the fill thread should call `usleep` at the end of each loop, after it has left its critical section. This call should be the only sleep call in your program.

transfer1 should take three command line arguments, the first being the name of the input file, the second being the name of the output file, the third being the sleep time.

Assuming this is the input file:

```
line 1
slightly longer
line 3
the longest line
line 5
```

Here is the output of an execution of transfer1 with a sleep time of 0:

```
$ ./transfer1 input.txt output.txt 0
fill thread: wrote [line 1
] into buffer
fill thread: wrote [slightly longer
] into buffer
fill thread: wrote [line 3
] into buffer
fill thread: wrote [the longest line
] into buffer
fill thread: wrote [line 5
] into buffer
fill thread: wrote QUIT into buffer
drain thread: read [line 1
] from buffer
drain thread: read [slightly longer
] from buffer
drain thread: read [line 3
] from buffer
drain thread: read [the longest line
] from buffer
drain thread: read [line 5
] from buffer
drain thread: read [QUIT] from buffer
```

Notice that the fill thread ran first and put all its lines into the buffer. Only then did the drain thread run, and it removed all lines. There was no concurrency in this execution. Each thread runs for only a short while and completes its task during that time, so the fill thread runs to completion.

Here is the output of an execution of transfer1 with a sleep time of 1000:

```
$ ./transfer1 input.txt output.txt 1000
fill thread: wrote [line 1
] into buffer
drain thread: read [line 1
] from buffer

(thousands of this line deleted:
drain thread: no new string in buffer)
```

3

```
fill thread: wrote [slightly longer
] into buffer
drain thread: read [slightly longer
] from buffer

(thousands of this line deleted:
drain thread: no new string in buffer)

fill thread: wrote [line 3
] into buffer
drain thread: read [line 3
] from buffer

(thousands of this line deleted:
drain thread: no new string in buffer)

fill thread: wrote [the longest line
] into buffer
drain thread: read [the longest line
] from buffer

(thousands of this line deleted:
drain thread: no new string in buffer)

fill thread: wrote [line 5
] into buffer
drain thread: read [line 5
] from buffer

(thousands of this line deleted:
drain thread: no new string in buffer)

fill thread: wrote QUIT into buffer
drain thread: read [QUIT] from buffer
```

The fill thread sleeping outside its critical section for a millisecond allowed the drain thread to be scheduled many times during that millisecond. However, each time the drain thread ran while the fill thread was sleeping there was no new line in the buffer. This is extremely wasteful: only 6 lines will ever be written into the buffer; the drain buffer ran needlessly thousands of times.

1.3 Part 3

Write a two-threaded program named "transfer2" that is the same as transfer1 except that it uses two semaphores for buffer synchronization as explained in lecture.

The signaling provided by the two semaphores eliminates all needless thread scheduling. Here is the output of an execution of transfer2 with a sleep time of 1000:

```
$ ./transfer2 input.txt output.txt 1000
fill thread: wrote [line 1
] into buffer
drain thread: read [line 1
] from buffer
fill thread: wrote [slightly longer
] into buffer
drain thread: read [slightly longer
] from buffer
fill thread: wrote [line 3
] into buffer
drain thread: read [line 3
] from buffer
fill thread: wrote [the longest line
] into buffer
drain thread: read [the longest line
] from buffer
fill thread: wrote [line 5
] into buffer
drain thread: read [line 5
] from buffer
fill thread: wrote QUIT into buffer
drain thread: read [QUIT] from buffer
```

There are no unnecessary thread schedulings. Each thread waits for the other thread's "post" operation.

2 Submission Instructions

Submit rw.c, transfer1.c, and transfer2.c via Moodle by the indicated date and time. Moodle will shut you out at the deadline. Since your clock and

Moodle's clock may differ by a few minutes, be sure to submit at least several minutes before the deadline.

Please note: *late work is not accepted; late submissions will not be graded and will receive a score of zero.* Also, points will be deducted if files have names different from `rw.c`, `transfer1.c`, and `transfer2.c`.

Your code should compile on `linux-lab.srcit.stevens.edu` with no errors or warnings using a compilation command such as:

```
gcc -Wall -pedantic-errors transfer1.c -o transfer1
```

Be sure to start early and contact the instructors if you have questions or encounter problems. Do NOT write the entire program then start debugging. Instead, write a part, debug it, then write the next part. Remove extra print statements from a section of code only once you're confident that the section is working.